Token Economics

By

Anand Kulkarni

anand.kulkarni1@zensar.com

Table of Content

Tokenization basics

- Tokenization is the process of breaking text into smaller units called tokens, which helps machines process and analyze language effectively.
- A token can be a word, a character, a subword fragment, or even an entire sentence, depending on the tokenizer being used.
- Every NLP model and every Large Language Model relies on tokenization as the first step. Without it, raw human language cannot be mapped into numerical vectors that AI can process.
- This is why tokenization is considered the foundation of Natural Language Processing and the core of modern AI systems.

Sample Tokenization program

```python
import tiktoken
encoding = tiktoken.get_encoding('cl100k_base')
text = 'Generative AI (Gen AI) is a branch of artificial intelligence focused on content creation'
tokens = encoding.encode(text)
print(tokens)
print(len(tokens))
print('Decoded text ', encoding.decode(tokens))

for token in tokens:
    print(f"{encoding.decode([token])} - {token}")
```

AI Tokens

- Tokens are small units of data — words, sub-words, characters, or pixel patches — that AI models process during training and inference.
- Roughly 1,000 tokens equal about 750 English words; for Gemini models, one token equals about four characters.
- Tokenization converts raw inputs (text, images, audio, video) into numerical sequences a model can learn from.
- AI providers bill per token, with output tokens costing 3–5× more than input tokens.
- Context windows — measured in tokens — set how much data a model can handle at once, from a few thousand to over one million tokens.

AI Tokens

- Reasoning models generate internal "thinking tokens" that multiply compute costs by 100× or more per prompt.
- Token prices are falling fast: median costs dropped roughly 200× per year between 2024 and 2026.
- Tokens are the atomic units of every generative AI system.
- Every prompt a user sends, every answer a model returns, and every image or audio clip it processes gets broken into these small, numbered fragments before the AI does any actual work.
- A token might be a full word like "cat," a sub-word like "ness," a single character, or a patch of pixels — depending on the data type and the tokenizer the model uses.

AI Tokens

- This makes tokens both the operating language and the billing unit of generative AI.
- Providers like OpenAI, Google, and Anthropic charge developers for every input and output token their APIs process.
- As of early 2026, lightweight models like Gemini Flash handle input at $0.08 per million tokens, while premium reasoning models like GPT-5.2 charge $14 per million output tokens. The gap between those two figures explains why understanding tokens is no longer optional for anyone building, buying, or scaling AI products.

How Models Use Tokens During Training?

- Training an AI model begins with tokenizing the full training dataset.
- For large language models, that dataset can contain billions or trillions of tokens.
- Generally, larger token counts during training lead to higher-quality models.
- The core training loop works through prediction.
- The model sees a sequence of tokens and tries to guess the next one.
- When it guesses wrong, internal parameters update to improve accuracy on the next attempt. This cycle repeats across the entire dataset until the model reaches a target accuracy threshold — a state called model convergence.

How Models Use Tokens During Training?

- After pretraining, a second phase called post-training narrows the model’s focus. Here, the model continues learning on a curated subset of tokens relevant to a specific domain or task - legal documents, medical records, business data, or conversational formats. The objective is to fine-tune the model, so it produces accurate, useful responses when deployed.

Tokens at Work: Inference and Reasoning

- Inference is where tokens meet the end user. A person submits a prompt — text, an image, an audio clip, a video - and the model translates that input into tokens. It processes those input tokens, generates output tokens as its response, then converts the result into the user's expected format. The input and output can even be in different modalities, such as an English text prompt that produces a Japanese translation, or a written description that generates an image.

Mode Context Window

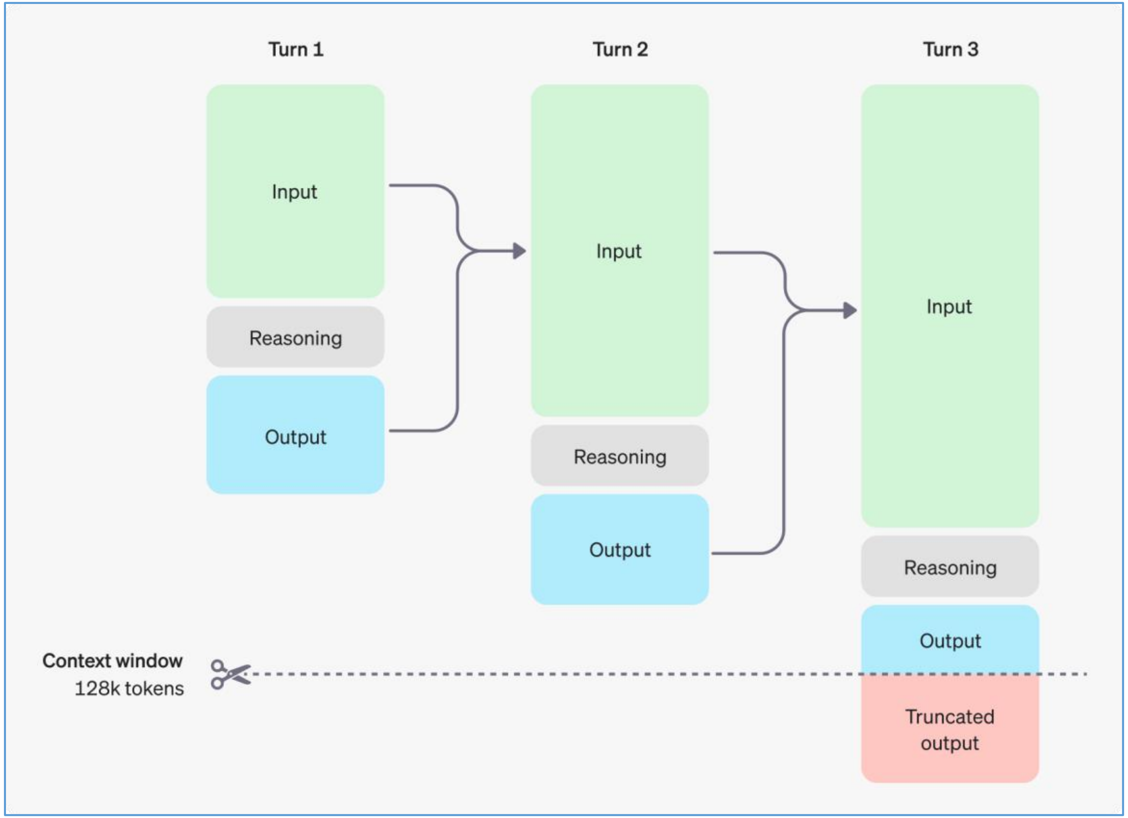

Mode Context Window

- Model's Context window is the maximum number of tokens it can process in a single interaction.
- A model with a few thousand tokens of context can handle a high-resolution image or a couple of pages of text.
- Models with context windows of tens of thousands of tokens can summarize a full novel or an hourlong podcast.
- Some models now offer context windows exceeding one million tokens, allowing users to feed in massive datasets for analysis in a single pass.
- For Gemini models, one token equals about four characters, and 100 tokens translate to roughly 60–80 English words. Developers can check context limits programmatically — Google's API exposes input and output token limits through its models endpoint, and a count_tokens function lets developers preview how many tokens a request will consume before sending it.

Reasoning Tokens

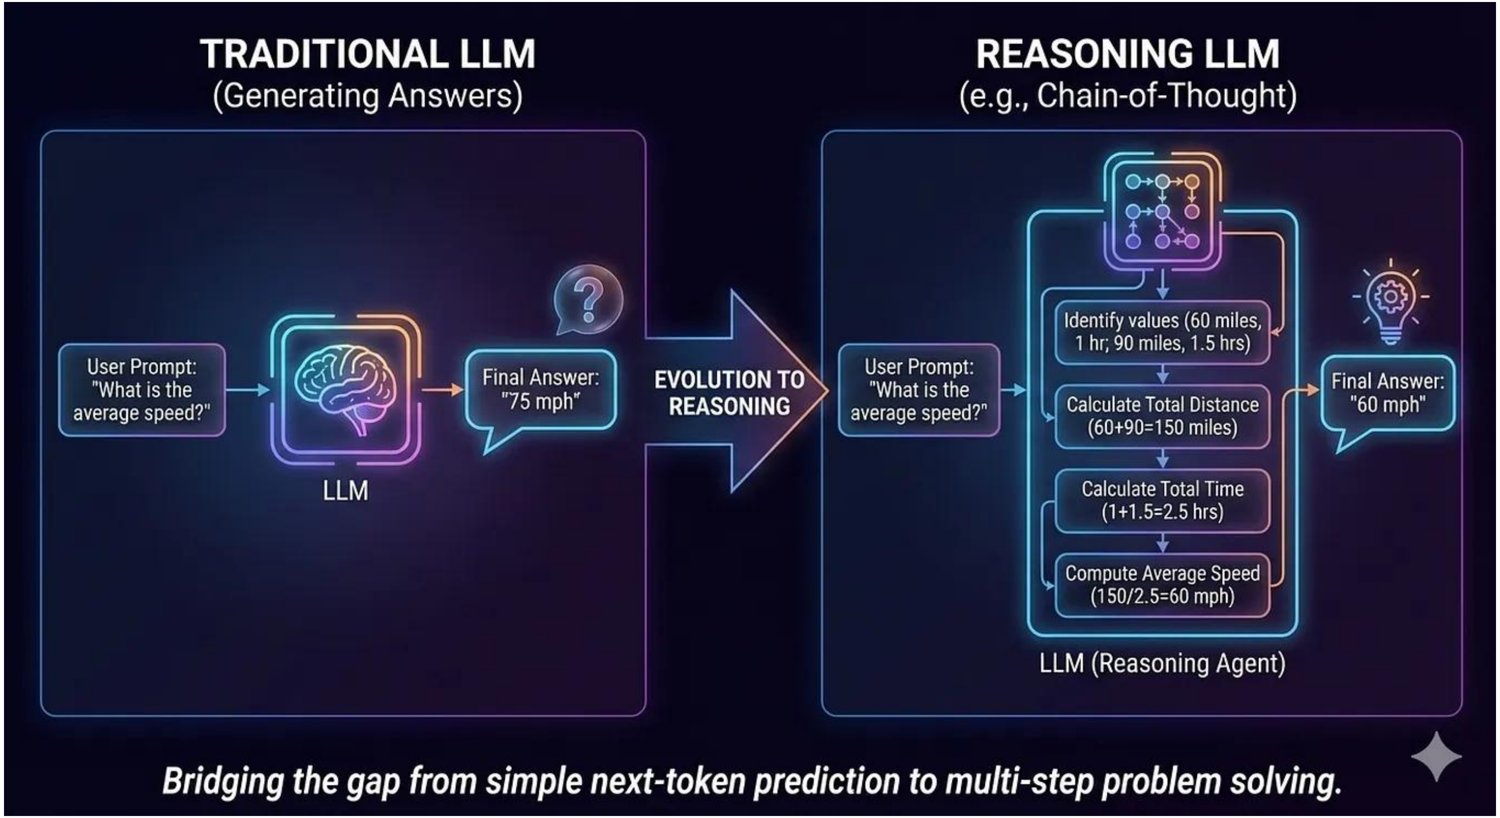

Reasoning Tokens

- In addition to input and output tokens, these models generate a large volume of internal "thinking tokens" as they work through multi-step problems.
- Thinking tokens are also called as Reasoning tokens.
- The reasoning tokens never appear in the user-facing response, but they occupy context-window space, consume compute, and get billed as output tokens.
- The impact on cost is significant. A single reasoning prompt can require 100× or more compute compared to a standard inference pass. This is sometimes called test-time scaling or long thinking — the model trades speed for quality, much like a human who takes longer to think through a harder question.

Multimodal Tokenization: Images, Video, Audio

- Generative AI is no longer text-only, and tokenization has adapted accordingly.
- Visual, audio, and video data all get converted into tokens at predictable rates.

Data Type	Tokenization Rate	Notes
Text	~1,000 tokens ≈ 750 English words	Varies by tokenizer and language
Small image (≤384px both sides)	258 tokens per image	Gemini API; fixed count
Large image (>384px)	258 tokens per 768×768 tile	Cropped and scaled into tiles
Video	263 tokens per second	Fixed rate; Gemini API
Audio	32 tokens per second	Fixed rate; Gemini API

How Tokens Drive AI Economics?

Model	Input (per 1M tokens)	Output (per 1M tokens)
Gemini 2.0 Flash Lite	$0.08	$0.30
GPT-4o Mini	$0.15	$0.60
GPT-4o	$2.50	$10.00
Claude 3.5 Sonnet	$3.00	$15.00
Claude Opus 4.6	$5.00	$25.00
GPT-5 (reasoning)	$15.00	$75.00

How Tokens Drive AI Economics?

- Tokens are the billing meter for the entire generative AI industry.
- Every API call consumes input tokens (the prompt) and output tokens (the response), each charged at per-million-token rates set by the provider.
- A consistent pattern across all providers: output tokens cost 3–5× more than input tokens because generating new text requires more computation than processing existing input.
- For a chatbot serving 10,000 daily users, even small differences in per-token pricing compound fast.
- One estimate puts the monthly cost for a mid-tier model at around $3,300 for 10 million tokens per day.

How Tokens Drive AI Economics?

- The token pricing has sharply come down.
- Between 2024 and 2026, median per-token costs dropped at a rate of roughly 200× per year.
- Prompt caching — where repeated context like system instructions gets processed at a reduced rate — adds another layer of savings. Anthropic's Claude Opus 4.6, for instance, charges $5 per million input tokens at standard rates but only $0.50 for cached reads, a 90% discount.
- Online token and its cost calculator: https://tokencalculator.app/

Token Limits and User Experience

- Beyond cost, tokens shape how AI services feel to use. Two performance metrics matter most - Time to first token and Inter-token latency.
- Time to first token measures the delay between a user submitting a prompt and the model beginning its response.
- Inter-token latency measures how fast subsequent tokens stream out after that.
- For chatbots, low time-to-first-token keeps conversations feeling natural.
- For text generation, optimized inter-token latency lets output match a person’s reading speed.
- Some providers set per-minute token ceilings for individual users to manage server load during high-demand periods.
- Others offer tiered plans where customers buy a fixed token budget shared between input and output.

What is AI Tokenomics?

- In AI, tokenomics refers to the economic framework governing how tokens—the smallest units of data processed by models—are generated, consumed, priced, and optimized.
- Tokenomics impacts cost efficiency, model selection, and workflow design. For example, summarizing a 15,000-token document with a high-end model like Claude 3 Opus can cost ~60× more than using a cheaper model like Claude 3 Haiku for the same task. This makes matching the right model to the right task critical.

What is AI Tokenomics?

Key Cost-Control Strategies:

- **Tiered Model Waterfall** – Start with the cheapest model, escalate only when needed.
- **Context Editing** – Summarize older conversation history to reduce token load.
- **Output Constraints** – Use parameters like max_output_tokens to limit verbosity.
- **Total Cost of Ownership (TCO)** – Factor in GPU costs, storage, monitoring, and orchestration—not just token API fees.

Bottom Line Mastering tokenomics means treating tokens as currency for computational reasoning. Businesses that strategically manage token flow—choosing the right models, optimizing context, and aligning spend with value—will gain a sustainable competitive advantage in AI deployment.

Pillars of AI Tokenomics

The pillars of AI Tokenomics are interconnected and deeply interdependent. They are crucial for enterprises planning or scaling AI deployments. Here's a brief overview of each pillar:

1) Supply: This pillar deals with the production and distribution of tokens, which are the fundamental units of work and currency in AI. It involves the planning, costing, and implementation of AI capacity to meet the demand for AI services.

2) Demand: This pillar focuses on the demand for AI services and the infrastructure requirements to support them. It involves forecasting and optimizing the use of AI resources to ensure efficient consumption and maximize value.

Pillars of AI Tokenomics

3) Utility: This pillar sets the ceiling on what can be charged for AI services. It determines the value of a token based on the intelligence it carries, the speed at which it is delivered, and the use case for which it can be utilized.

4) Value: This pillar connects the total cost of AI to its outputs, including labor. It involves the allocation, forecasting, optimization, and cost to serve of AI workloads, expressed at cost per call rather than cost per token.

Thank you!!